



Chapter X: Big Data

Database System Concepts, 6th Ed.

©Silberschatz, Korth and Sudarshan

See www.db-book.com for conditions on re-use



Chapter X: Big Data

- Map Reduce
 - The Map reduce paradigm
 - Distributed File Systems
 - Hadoop
- Big Data Storage Systems



The MapReduce Paradigm

- Platform for reliable, scalable parallel computing
- Abstracts issues of distributed and parallel environment from programmer.
- Paradigm dates back many decades
 - But very large scale implementations running on clusters with 10^3 to 10^4 machines are more recent
 - Google Map Reduce, Hadoop, ..
- Data access done using distributed file systems



Distributed File Systems

- Highly scalable distributed file system for large data-intensive applications.
 - E.g. 10K nodes, 100 million files, 10 PB
- Provides redundant storage of massive amounts of data on cheap and unreliable computers
 - Files are replicated to handle hardware failure
 - Detect failures and recovers from them
- Examples:
 - Google File System (GFS)
 - Hadoop File System (HDFS)



MapReduce: File Access Count Example

- Given log file in following format:

...

2013/02/21 10:31:22.00EST [/slide-dir/11.ppt](#)

2013/02/21 10:43:12.00EST [/slide-dir/12.ppt](#)

2013/02/22 18:26:45.00EST [/slide-dir/13.ppt](#)

2013/02/22 20:53:29.00EST [/slide-dir/12.ppt](#)

...

- Goal: find how many times each of the files in the slide-dir directory was accessed between 2013/01/01 and 2013/01/31.
- Options:
 - Sequential program too slow on massive datasets
 - Load into database expensive, direct operation on log files cheaper
 - Custom built parallel program for this task possible, but very laborious
 - Map-reduce paradigm



MapReduce Programming Model

- Inspired from map and reduce operations commonly used in functional programming languages like Lisp.
- Input: a set of key/value pairs
- User supplies two functions:
 - $\text{map}(k,v) \rightarrow \text{list}(k1,v1)$
 - $\text{reduce}(k1, \text{list}(v1)) \rightarrow v2$
- $(k1,v1)$ is an intermediate key/value pair
- Output is the set of $(k1,v2)$ pairs
- For our example, assume that system
 - breaks up files into lines, and
 - calls map function with value of each line
 - ▶ Key is the line number

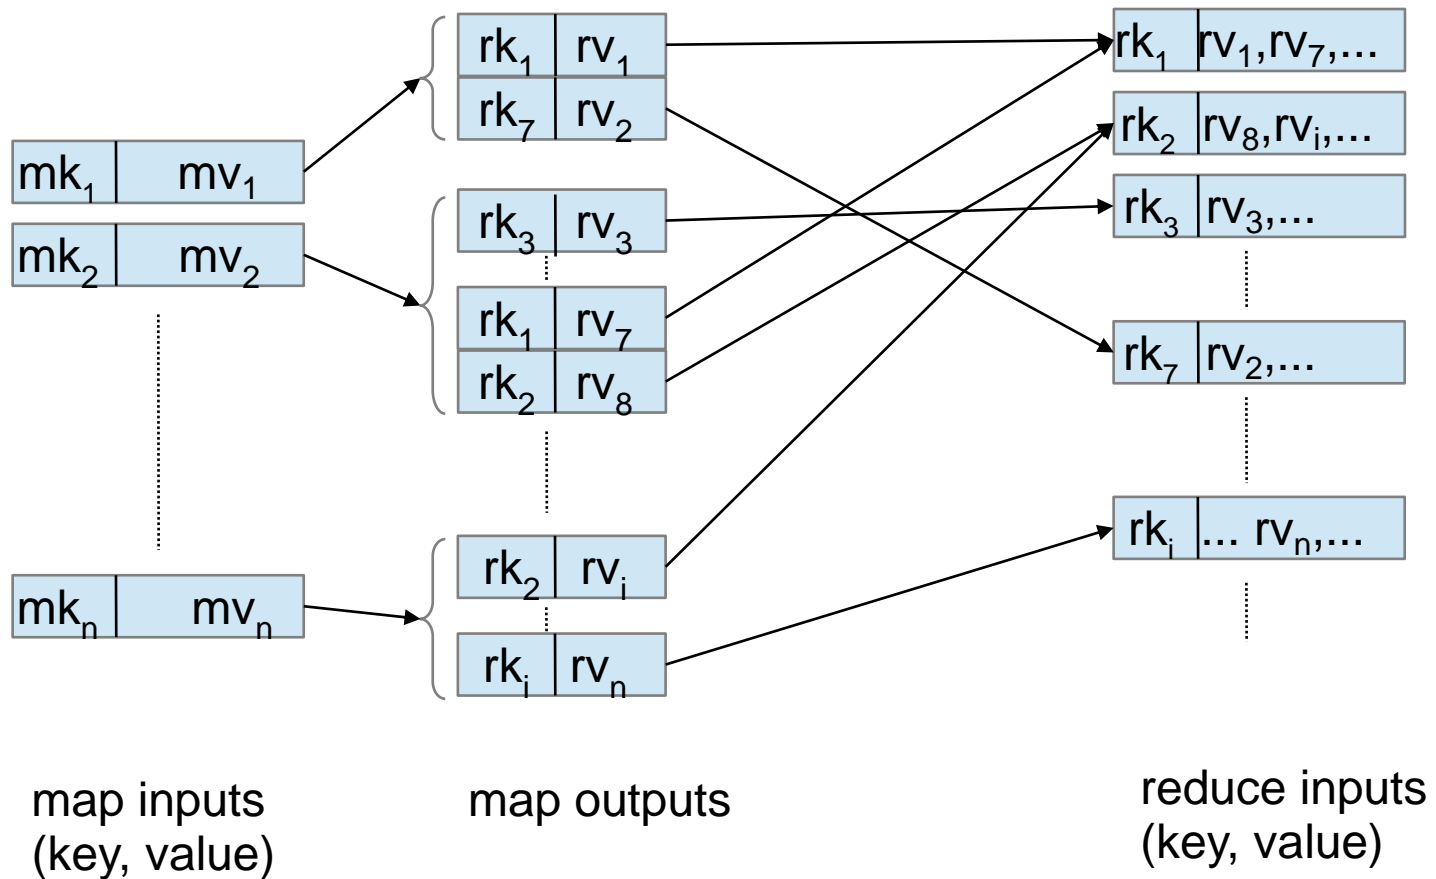


MapReduce: File Access Count Example

```
map(String key, String record) {  
    String attribute[3];  
    .... break up record into tokens (based on space character), and store the  
           tokens in array attributes  
    String date = attribute[0];  
    String time = attribute[1];  
    String filename = attribute[2];  
    if (date between 2013/01/01 and 2013/01/31  
        and filename starts with "/slide-dir/")  
        emit(filename, 1).  
}  
  
reduce(String key, List recordlist) {  
    String filename = key;  
    int count = 0;  
    For each record in recordlist  
        count = count + 1.  
    output(filename, count)  
}
```



Schematic Flow of Keys and Values



■ Flow of keys and values in a map reduce task



MapReduce: Word Count Example

- Consider the problem of counting the number of occurrences of each word in a large collection of documents
- How would you do it in parallel ?
- Solution:
 - Divide documents among workers
 - Each worker parses document to find all words, map function outputs (word, count) pairs
 - Partition (word, count) pairs across workers based on word
 - For each word at a worker, reduce function locally add up counts
- Given input: “One a penny, two a penny, hot cross buns.”
 - Records output by the map() function would be
 - ▶ (“One”, 1), (“a”, 1), (“penny”, 1), (“two”, 1), (“a”, 1), (“penny”, 1), (“hot”, 1), (“cross”, 1), (“buns”, 1).
 - Records output by reduce function would be
 - ▶ (“One”, 1), (“a”, 2), (“penny”, 2), (“two”, 1), (“hot”, 1), (“cross”, 1), (“buns”, 1)



Pseudo-code

map(String input_key, String input_value):

// input_key: document name

// input_value: document contents

for each word w in input_value:

Emit(w, "1");

// Group by step done by system on key of intermediate Emit above,

// and reduce called on list of values in each group.

reduce(String output_key, Iterator intermediate_values):

// output_key: a word

// output_values: a list of counts

int result = 0;

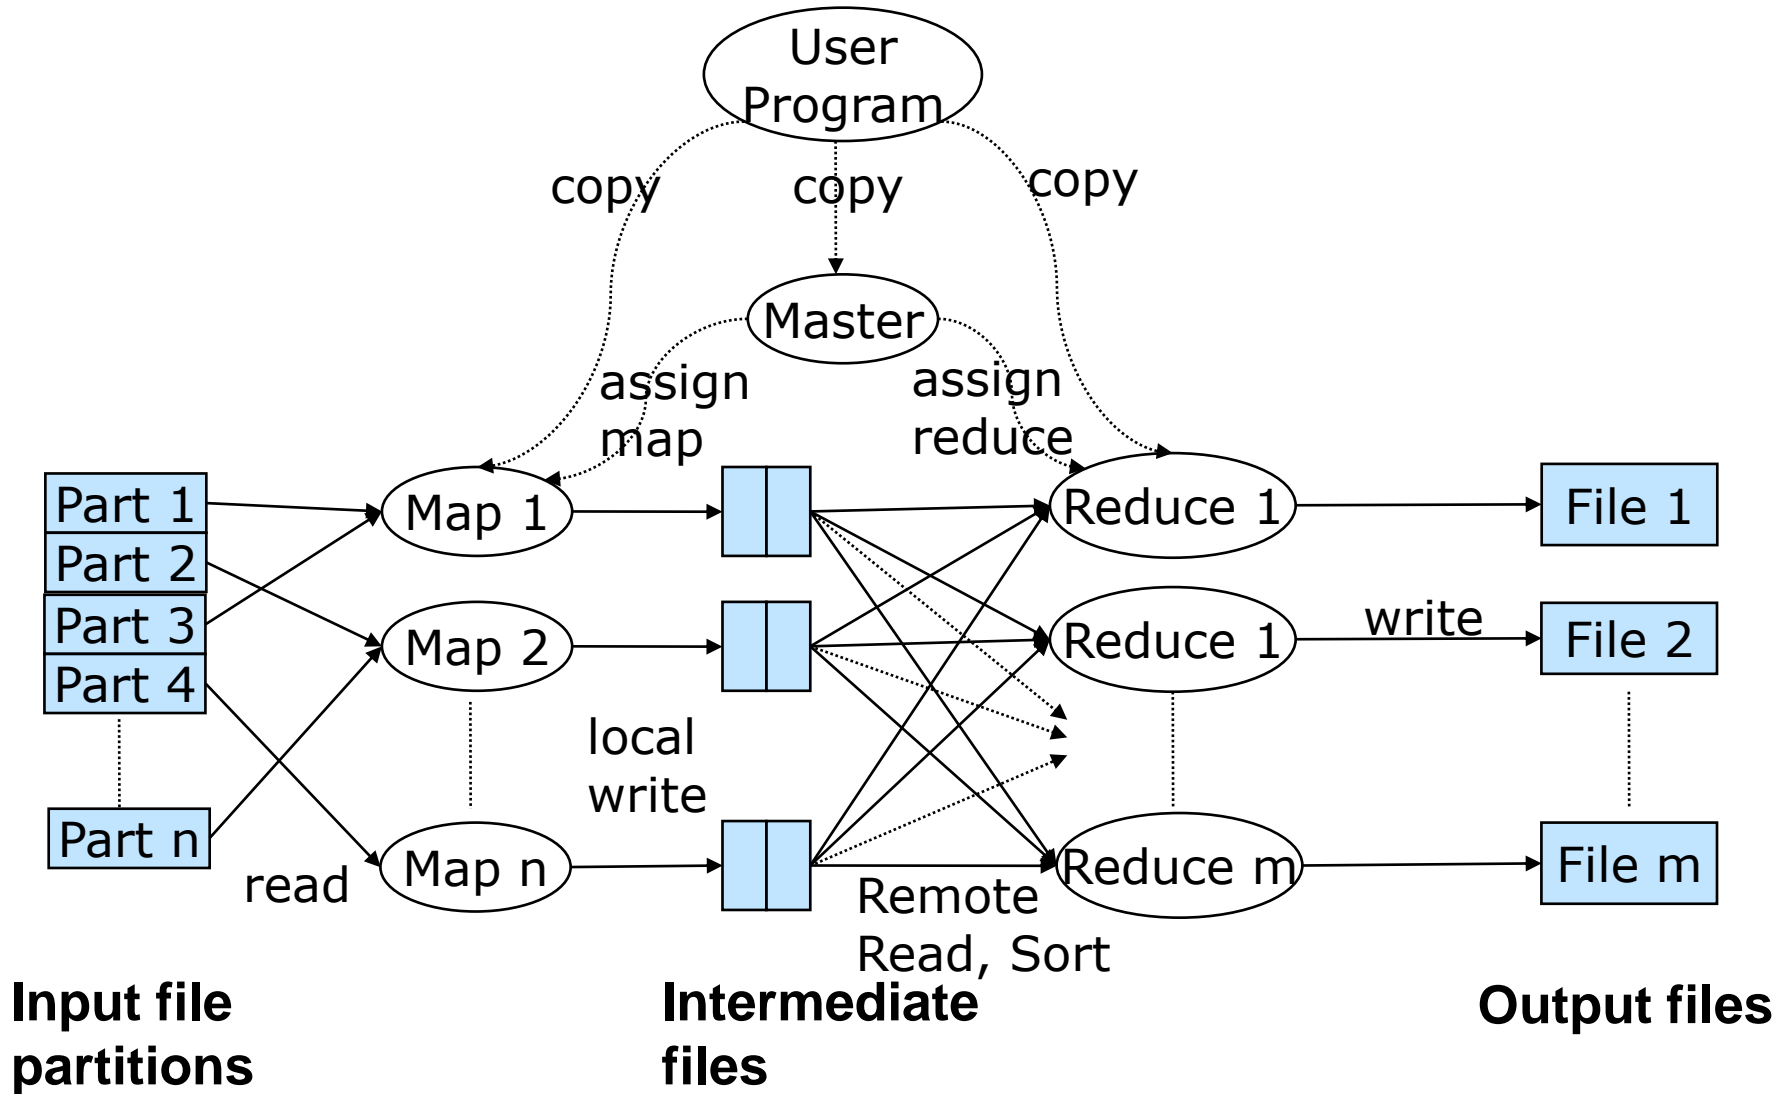
for each v in intermediate_values:

result += ParseInt(v);

Output(result);



Parallel Processing of MapReduce Job





Hadoop

- Google pioneered map-reduce implementations that could run on thousands of machines (nodes), and transparently handle failures of machines
- Hadoop is a widely used open source implementation of Map Reduce written in Java
 - Map and reduce functions can be written in several different languages, we use Java.
- Input and output to map reduce systems such as Hadoop must be done in parallel
 - Google used GFS distributed file system
 - Hadoop uses Hadoop File System (HDFS)
 - ▶ File blocks partitioned across many machines
 - ▶ Blocks are replicated so data is not lost/unavailable if a machine crashes
 - ▶ Central “name node” provides metadata such as which blocks are contained in which files



Hadoop

■ Types in Hadoop

- Generic Mapper and Reducer interfaces both take four type arguments, that specify the types of the
 - ▶ input key, input value, output key and output value
- Map class in next slide implements the Mapper interface
 - ▶ Map input key is of type LongWritable, i.e. a long integer
 - ▶ Map input value which is (all or part of) a document, is of type Text.
 - ▶ Map output key is of type Text, since the key is a word,
 - ▶ Map output value is of type IntWritable, which is an integer value.



Hadoop Code in Java: Map Function

```
public static class Map extends Mapper<LongWritable, Text, Text, IntWritable>
{
    private final static IntWritable one = new IntWritable(1);
    private Text word = new Text();
    public void map(LongWritable key, Text value, Context context)
        throws IOException, InterruptedException
    {
        String line = value.toString();
        StringTokenizer tokenizer = new StringTokenizer(line);
        while (tokenizer.hasMoreTokens()) {
            word.set(tokenizer.nextToken());
            context.write(word, one);
        }
    }
}
```



Hadoop Code in Java: Reduce Function

```
public static class Reduce extends Reducer<Text, IntWritable, Text,
    IntWritable> {
    public void reduce(Text key, Iterable<IntWritable> values,
        Context context) throws IOException, InterruptedException
    {
        int sum = 0;
        for (IntWritable val : values) {
            sum += val.get();
        }
        context.write(key, new IntWritable(sum));
    }
}
```



Hadoop Job Parameters

- The classes that contain the map and reduce functions for the job
 - set by methods `setMapperClass()` and `setReducerClass()`
- The types of the job's output key and values
 - set by methods `setOutputKeyClass()` and `setOutputValueClass()`
- The input format of the job
 - set by method `job.setInputFormatClass()`
 - ▶ Default input format in Hadoop is the `TextInputFormat`,
 - map key whose value is a byte offset into the file, and
 - map value is the contents of one line of the file
- The directories where the input files are stored, and where the output files must be created
 - set by `addInputPath()` and `addOutputPath()`
- And many more parameters



Hadoop Code in Java: Overall Program

```
public class WordCount {  
    public static void main(String[] args) throws Exception {  
        Configuration conf = new Configuration();  
        Job job = new Job(conf, "wordcount");  
        job.setOutputKeyClass(Text.class);  
        job.setOutputValueClass(IntWritable.class);  
        job.setMapperClass(Map.class);  
        job.setReducerClass(Reduce.class);  
        job.setInputFormatClass(TextInputFormat.class);  
        job.setOutputFormatClass(TextOutputFormat.class);  
        FileInputFormat.addInputPath(job, new Path(args[0]));  
        FileOutputFormat.setOutputPath(job, new Path(args[1]));  
        job.waitForCompletion(true);  
    }  
}
```



Local Pre-Aggregation

- Combiners: perform partial aggregation to minimize network traffic
 - E.g. within machine
 - And/or at rack level
- In Hadoop, reduce function is used by default if combiners are enabled
 - But alternative implementation of combiner can be specified if input and output types of reducers are different



Implementations

- Google
 - Not available outside Google
- Hadoop
 - An open-source implementation in Java
 - Uses HDFS for stable storage
 - Download: <http://lucene.apache.org/hadoop/>
- Aster Data
 - Cluster-optimized SQL Database that also implements MapReduce
 - ▶ IITB alumnus among founders
- And several others, such as Cassandra at Facebook, etc.



Map Reduce vs. Parallel Databases

- Map Reduce widely used for parallel processing
 - Google, Yahoo, and 100' s of other companies
 - Example uses: compute PageRank, build keyword indices, do data analysis of web click logs,
- Database people say: but parallel databases have been doing this for decades
- Map Reduce people say:
 - we operate at scales of 1000' s of machines
 - We handle failures seamlessly
 - We allow procedural code in map and reduce and allow data of any type
 - ▶ many real-world uses of MapReduce that cannot be expressed in SQL.



Map Reduce vs. Parallel Databases (Cont.)

- Map Reduce is cumbersome for writing simple queries
- Current approach: declarative querying, with execution on Map Reduce infrastructure
 - Pig Latin – Declarative language supporting complex data using JSON; From Yahoo
 - ▶ Programmer has to specify parser for each input source
 - Hive – SQL based syntax; From Facebook
 - ▶ Allows specification of schema for each input source
 - ▶ Has become very popular
 - SCOPE system from Microsoft
- Many proposed extensions of Map Reduce to allow joins, pipelining of data, etc.



Big Data Storage Systems

- Need to store massive amounts of data
- Scalability: ability to grow the system by adding more machines
 - Scalability of storage volume, and access rate
- Distributed file systems good for scalable storage of unstructured data
 - E.g. log files
- Massively parallel database ideal for storing records
 - But hard to support all database features across thousands of machines
- Massively parallel key-value stores built to support scalable storage and access of records
 - E.g. Big Table from Google, PNUTS/Sherpa from Yahoo, HBase from Apache Hadoop project
- Applications using key-value stores manage query processing on their own




Key Value Stores

- Key-value stores support
 - `put(key, value)`: used to store values with an associated key,
 - `get(key)`: which retrieves the stored value associated with the specified key.
- Some systems such as Bigtable additionally provide range queries on key values
- Multiple versions of data may be stored, by adding a timestamp to the key



PNUTS Parallel Storage System Architecture





Data Representation

- Records in many big data applications need to have a flexible schema
 - Not all records have same structure
 - Some attributes may have complex substructure
- XML and JSON data representation formats widely used
- An example of a JSON object is:

```
{  
  "ID": "22222",  
  "name": {  
    "firstname": "Albert",  
    "lastname": "Einstein"  
  },  
  "deptname": "Physics",  
  "children": [  
    { "firstname": "Hans", "lastname": "Einstein" },  
    { "firstname": "Eduard", "lastname": "Einstein" }  
  ]  
}
```